


Hive



Agenda

- What is Hive?
 - Java vs Hive
 - Components
 - Configuration
- Hive Queries
 - Tables
 - Explain
 - Join
- Hive and HBase
 - HBase Table Mapping
 - Multiple columns to MAP



SQL for Hadoop

- **Data warehouse augmentation is a very common use case for Hadoop**
- **While highly scalable, MapReduce is notoriously difficult to use**
 - Java API is tedious and requires programming expertise
 - Unfamiliar languages (e.g. Pig) also requiring expertise
 - Many different file formats, storage mechanisms, configuration options, etc.
- **SQL support opens the data to a much wider audience**
 - Familiar, widely known syntax
 - Common catalog for identifying data and structure
 - Clear separation of defining the *what* (you want) vs. the *how* (to get it)

What is Hive?

- **A system for managing and querying structured data built on top of Hadoop**
 - Map-Reduce for execution
 - HDFS for storage
 - Metadata on raw files
- **Key Building Principles:**
 - SQL is a familiar data warehousing language
 - Extensibility – Types, Functions, Formats, Scripts
 - Scalability and Performance

Java versus Hive: The Word Count Algorithm

```
package org.myorg;

import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class WordCount {

    public static class Map extends Mapper<LongWritable, Text, Text,
    IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException {
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

```
public static class Reduce extends Reducer<Text, IntWritable, Text,
IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values, Context
context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    Job job = new Job(conf, "wordcount");

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    job.setMapperClass(Map.class);
    job.setReducerClass(Reduce.class);

    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

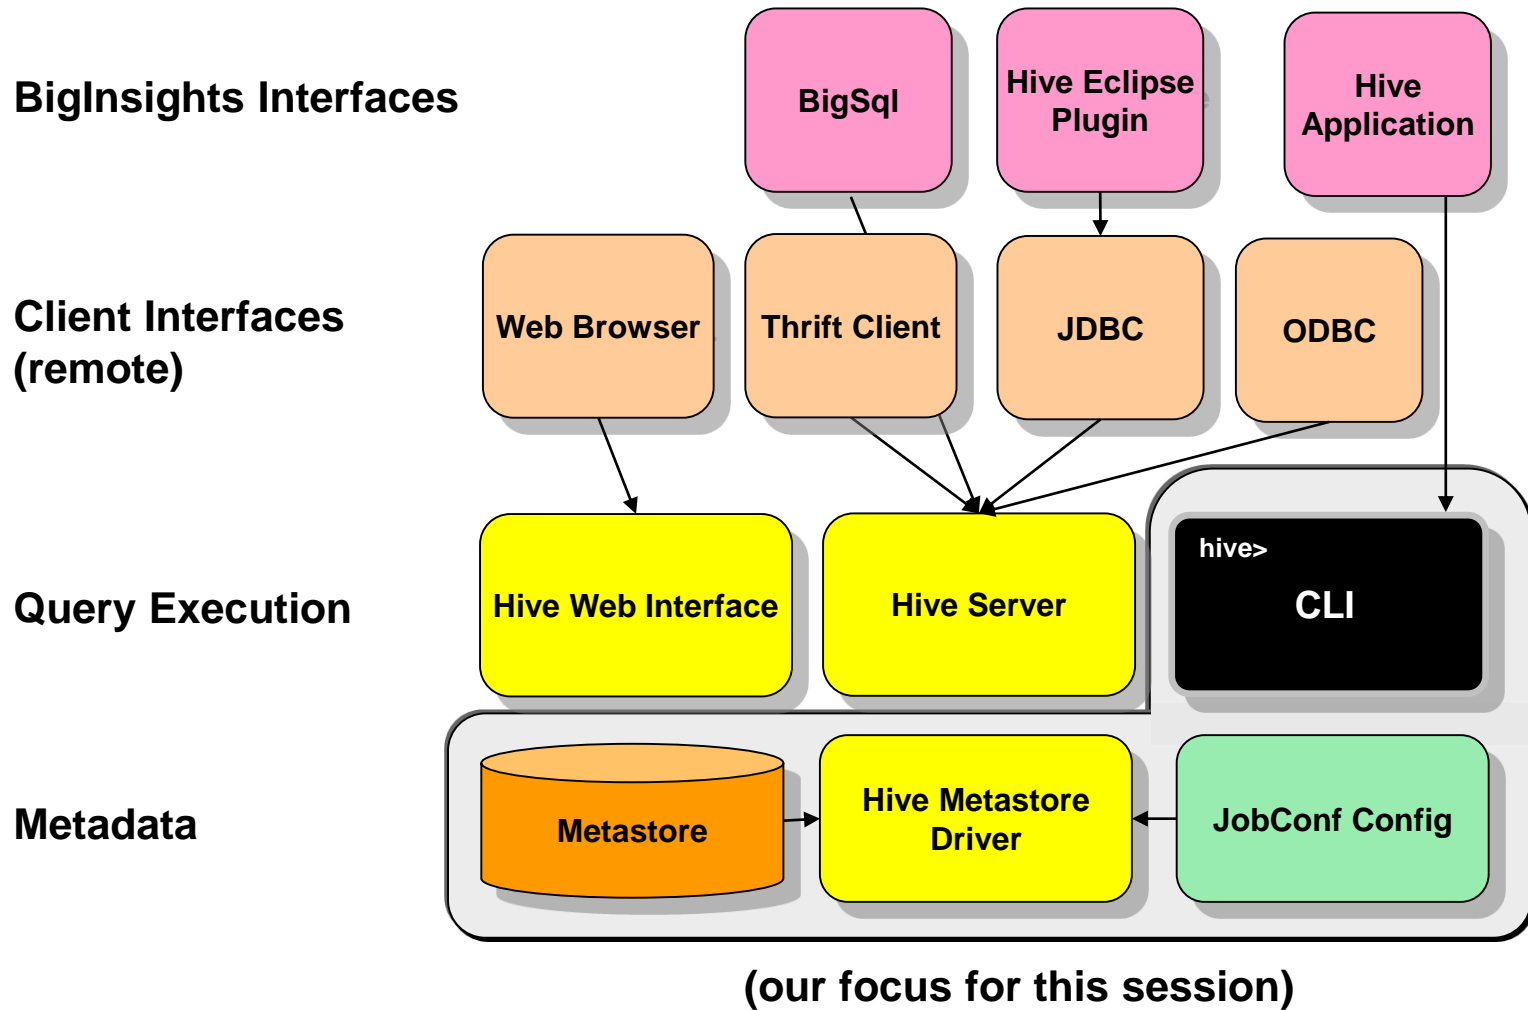
    job.waitForCompletion(true);
}
```

Hive and Word Count!

```
CREATE TABLE docs (line STRING);  
  
LOAD DATA INPATH 'docs' OVERWRITE INTO TABLE docs;  
  
CREATE TABLE word_counts AS  
SELECT word, count(1) AS count FROM  
      (SELECT explode(split(line, 's')) AS word FROM docs)  
GROUP BY word  
ORDER BY word;
```

*See “Programming Hive”, *Capriolo, Wampler & Rutherglen*

Hive - Components



Starting Hive – The Hive Shell

- **The Hive shell is located in**
`$HIVE_HOME/bin/hive`
- **From the shell you can**
 - Perform queries, DML, and DDL
 - View and manipulate table metadata
 - Retrieve query explain plans (execution strategy)

```
$ $HIVE_HOME/bin/hive
2013-01-14 23:36:52.153 GMT : Connection obtained for host: master-
Logging initialized using configuration in file:/opt/ibm/biginsight
Hive history file=/var/ibm/biginsights/hive/query/biadmin/hive_job

hive> show tables;
mytab1
mytab2
mytab3
OK
Time taken: 2.987 seconds
hive> quit;
```


Data Types and Models

- **Supports a number of scalar and structured data types**

- tinyint, smallint, int, bigint, float, double
- boolean
- string, binary
- timestamp
- array – e.g. `array<int>`
- struct – e.g. `struct<f1:int, f2:array<string>>`
- map – e.g. `map<int, string>`
- union – e.g. `uniontype<int, string, double>`

- **Partitioning**

- Can partition on one or more columns
- Value partitioning only, range partitioning is not yet supported

- **Bucketing**

- Sub-partitioning/grouping of data by hash within partitions
- Useful for sampling and improves some join operations

Data Model - Partition

- **Value partition based on partition columns**
- **Nested sub-directories in HDFS for each combination of partition column values**
- **Example**
 - Partition columns : ds, ctry
 - HDFS subdirectory for ds = 20090801, ctry = US
 - .../hive/warehouse/pview/ds=20090801/ctry=US
 - HDFS subdirectory for ds = 20090801, ctry = CA
 - .../hive/warehouse/pview/ds=20090801/ctry=CA

Data Model - Bucket

- **Split data based on hash of a column – mainly for parallelism**
- **One HDFS file per bucket within partition sub-directory**
- **Example**
 - Bucket column : user into 32 buckets
 - HDFS file for user hash 0
 - .../hive/warehouse/pview/ds=20090801/ctry=US/part-00000
 - HDFS file for user hash bucket 20
 - .../hive/warehouse/pview/ds=20090801/ctry=CA/part-00020

Data Model – External Table

- Point to existing data directories in HDFS
- Can create tables and partitions – partition columns just become annotations to external directories
- Example : create external table with partition

```
CREATE EXTERNAL TABLE pview (userid int, pageid int, ds string, ctry string)  
PARTITIONED ON (ds string, ctry string)  
STORED AS textfile  
LOCATION '/path/to/existing/table'
```

- Example : add a partition to external table

```
ALTER TABLE pview  
ADD PARTITION (ds='20090801', ctry='US')  
LOCATION '/path/to/existing/partition'
```

Data Model – External Table

- **For External Tables -**
- External table stores files on the HDFS server but tables are not linked to the source file completely.
- If you delete an external table the file still remains on the HDFS server.
- As an example if you create an **external table** called “**table_test**” in HIVE using HIVE-QL and link the table to file “**file**”, ***then deleting “table_test” from HIVE will not delete “file” from HDFS.***
- External table files are accessible to anyone who has access to HDFS file structure and therefore security needs to be managed at the HDFS file/folder level.
- **Meta data is maintained on master node, and deleting an external table from HIVE only deletes the metadata not the data/file.**
- **For Internal Tables-**
- Stored in a directory based on settings in hive.metastore.warehouse.dir, ***by default internal tables are stored*** in the following directory “**/user/hive/warehouse**” you can change it by updating the location in the config file .
- Deleting the table deletes the metadata and data from master-node and HDFS respectively.
- Internal table file security is controlled solely via HIVE. Security needs to be managed within HIVE, probably at the schema level (depends on organization)

Physical Layout

■ Hive warehouse directory structure

Base directory in HDFS for all Hive tables (BigInsights installed default)

/biginsights/hive/warehouse/

db1.db/

Database (schemas) are contained in ".db" subdirectories

tab1/

tab1.dat

Table "tab1"

part_tab1/
state=NJ/

part_tab1.dat

state=CA/

part_tab1.dat

Table partitioned by "state" column. One subdirectory per unique value. Query predicates eliminate partitions (directories) that need to be read

■ Data files are just regular HDFS files

- Internal format can vary table-to-table (delimited, sequence, etc.)

■ Supports "external" tables

Creating a table

■ Creating a delimited table

```
hive> create table users
(
  id          int,
  office_id   int,
  name        string,
  children    array<string>
)
row format delimited
  fields terminated by '|'
  collection items terminated by ':'
stored as textfile;
```

file: users.dat

```
1|1|Bob Smith|Mary
2|1|Frank Barney|James:Liz:Karen
3|2|Ellen Lacy|Randy:Martin
4|3|Jake Gray|
5|4|Sally Fields|John:Fred:Sue:Hank:Robert
```

■ Inspecting tables:

```
hive> show tables;
OK
users
Time taken: 2.542 seconds
```

```
hive> describe users;
OK
id          int
office_id   int
name        string
children    array<string>
Time taken: 0.129 seconds
```

Table population and querying

- **Loading data from input file**

```
hive> load data local inpath 'users.dat' into table users;  
Copying data from file:/home/biadmin/hive_demo/users.dat  
Copying file: file:/home/biadmin/hive_demo/users.dat  
Loading data to table default.users  
OK  
Time taken: 0.276 seconds
```

- **The “local” indicates the source data is on the local (unix) filesystem**

- Otherwise file is assumed to be on HDFS

- **Our first query:**

```
hive> select * from users;  
1      1      Bob Smith      ["Mary"]  
2      1      Frank Barney   ["James", "Liz", "Karen"]  
3      2      Ellen Lacy     ["Randy", "Martin"]  
4      3      Jake Gray      []  
5      4      Sally Fields   ["John", "Fred", "Sue", "Hank", ...]
```


Tables derived from queries

- **Tables may be created using queries on other tables**

```
create table emps_by_state
as
select o.state as state, count(*) as employees
  from office o left outer join users  u
    on u.office_id = o.office_id
group by o.state;
```

- **or...**

```
create table emps_by_state ... stored as textfile;

insert overwrite table emps_by_state
select o.state as state, count(*) as employees
  from office o left outer join users  u
    on u.office_id = o.office_id
group by o.state;
```

Storage file formats

- **The STORED AS clause indicates the storage file/record format on HDFS**

- TEXTFILE – Stored as a text line (one line per record)
- SEQUENCEFILE – Stored as a Hadoop sequence file
- RCFILE – Semi-columnar data storage with good compression. Best performance of the built-in storage formats
- INPUTFORMAT/OUTPUTFORMAT – Can provide a specific Hadoop input or output format class

```
create table foo ( ... )  
row format delimited fields terminated by ','  
storage format sequencefile;
```

- **This just dictates how *records* (rows) are stored but not values**
 - RCFILE is an exception, it has its own value storage formats

Record Format

- **Values are written/read into file using a Hive SerDe**
 - **Serializer/Deserializer** class – encodes and decodes values
- **Default SerDe is Hive's LazySimpleSerDe**
 - Delimited format
 - Delimiters specified with `DELIMITED` clause
 - “Lazy” indicates values are only read if accessed in the query
 - Re-uses objects across rows
 - `TIMESTAMP` must be formatted as `yyyy-mm-dd hh:mm:ss.ffffff`
 - Null value can be provided as `\N`
- **SerDe defined with the `STORED BY` clause**

```
create table foo ( ... )  
storage format sequencefile  
stored by 'org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe'
```

Record Format (cont.)

- **Other built-in SerDe implementations**
 - Thrift (ThriftSerDe) – Thrift binary encoding
 - Regular Expressions (RegexSerDe) – Use regex to extract fields
 - Hive Binary format (LazyBinarySerDe) – Binary storage
- **You can implement your own SerDe**
- **Others available freely on the internet**
 - JSON
 - Avro
 - Etc.

Explain

- The **explain** keyword generates a Hive explain plan for the query

```
hive> explain select o.state as state, count(*) as employees
      from office o left outer join users u
      on u.office_id = o.office_id
group by o.state
order by state;
```

```
STAGE DEPENDENCIES:
  Stage-1 is a root stage
  Stage-2 depends on stages: Stage-1
  Stage-3 depends on stages: Stage-2
  Stage-0 is a root stage
...
```

- Query requires three MapReduce jobs to implement
 - One to join the two inputs
 - One to perform the GROUP BY
 - One to perform the final sort

Joins

- Hive supports joins via ANSI join syntax only




```
select ...  
  from T1, T2  
 where T1.a = T2.b
```



```
select ...  
  from T1 join T2  
       on T1.a = T2.b
```

- Example join

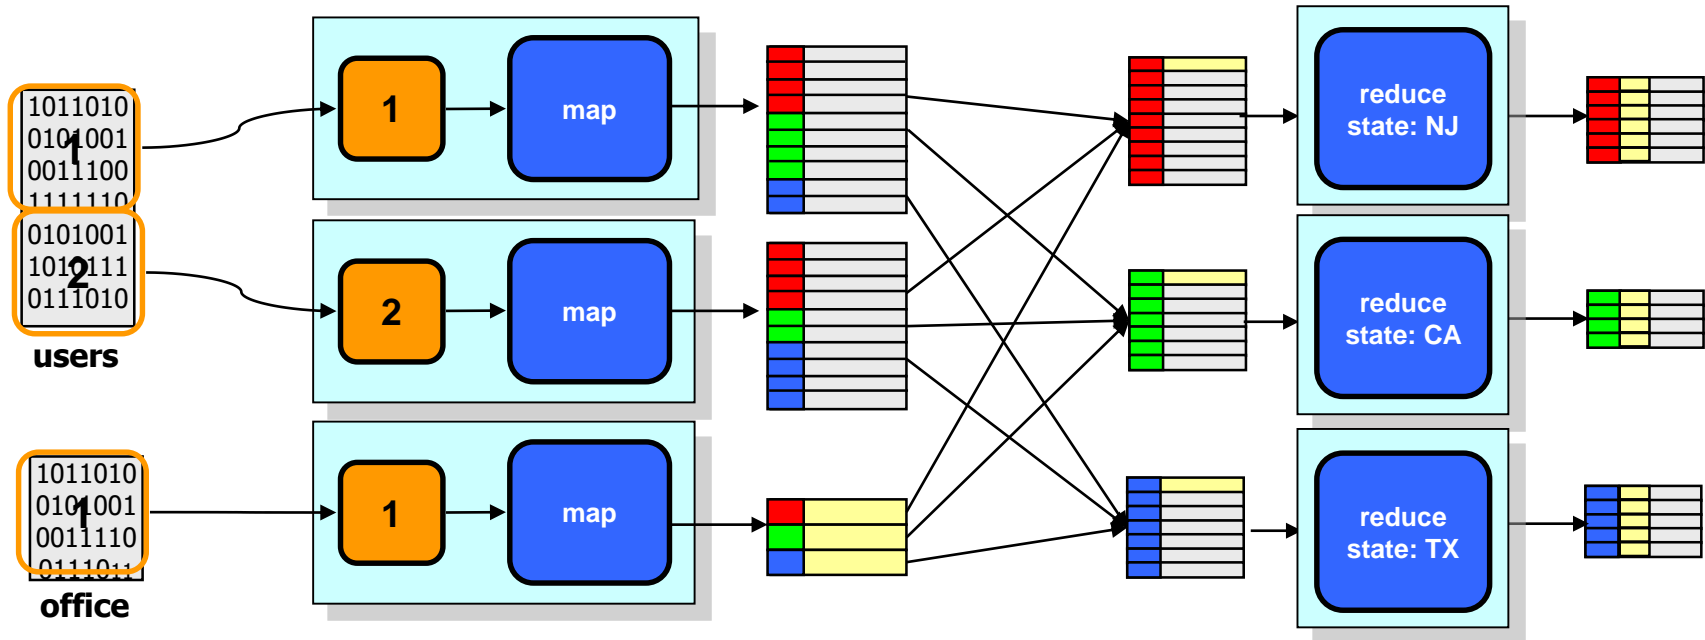
```
hive> select o.state as state, count(*) as employees  
       from office o left outer join users u  
         on u.office_id = o.office_id  
       group by o.state  
       order by state;
```



- On my (crummy) demo environment this takes 67 seconds (!)
 - Only five offices and five users (< 1KB of data)
 - What's going on?

Understanding distributed joins

- Not entirely necessary, but very useful to understand:



- **Hive only supports equi-joins ($t1.c1 = t2.c2$)**
 - E.g. it cannot perform non-equi-joins - $t1.c1 > t2.c2$
 - From the above diagram, do you understand why?

User Defined Function

▪ Java Code

```
package com.example.hive.udf;

import org.apache.hadoop.hive.ql.exec.UDF;
import org.apache.hadoop.io.Text;

public final class Lower extends UDF {
    public Text evaluate(final Text s) {
        if (s == null) { return null; }
        return new Text(s.toString().toLowerCase());
    }
}
```

▪ Registering the Class

- **CREATE FUNCTION** my_lower **AS** 'com.example.hive.udf.Lower';

▪ Using the Function

- **SELECT** my_lower(title), sum(freq) **FROM** titles **GROUP BY** my_lower(title);

Hive and HBase

- **Hive comes with an HBase *storage handler***
- **Allows MapReduce queries and loading of HBase tables**
- **Uses predicate pushdown to optimize query**
 - Scans only necessary regions based upon table key
 - Applies predicates as HBase row filters (if possible)
- **Usually Hive must be provided additional jars and configuration in order to work with HBase**

```
$ hive \  
  --auxpath \  
    $HIVE_SRC/build/dist/lib/hive-hbase-handler-0.9.0.jar,\  
    $HIVE_SRC/build/dist/lib/hbase-0.92.0.jar,\  
    $HIVE_SRC/build/dist/lib/zookeeper-3.3.4.jar,\  
    $HIVE_SRC/build/dist/lib/guava-r09.jar \  
  -hiveconf hbase.master=hbase.yoyodyne.com:60000
```

- **BigInsights' Hive is preconfigured for HBase integration, so this is unnecessary**

HBase table mapping

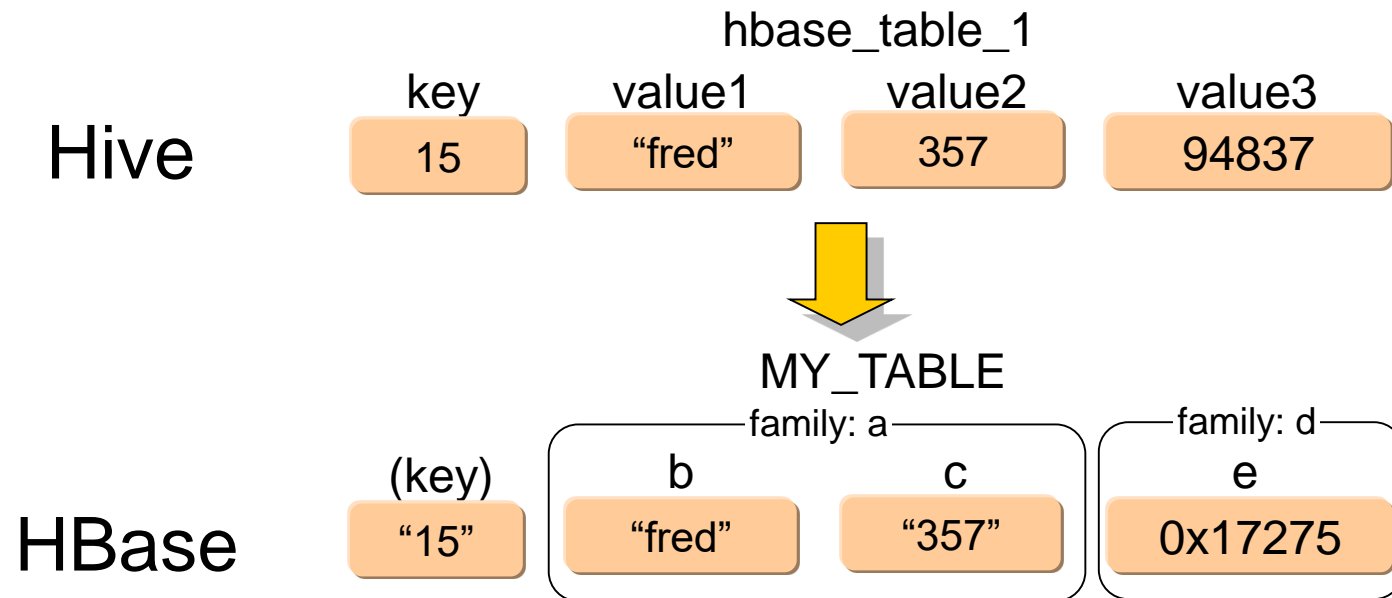
■ Creating an HBase table

```
CREATE TABLE hbase_table_1 (  
    key      int,  
    value1 string,  
    value2 int,  
    value3 int)  
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler`'  
WITH SERDEPROPERTIES (  
    "hbase.columns.mapping" = ":key,a:b,a:c,d:e"  
)  
TBLPROPERTIES(  
    "hbase.table.name" = "MY_TABLE"  
); ;
```

- **The `hbase.table.name` property provides the table name in HBase**
 - Optional. If not provided, the Hive name is assume

HBase table mapping

```
CREATE TABLE hbase_table_1 (  
  key      int,  
  value1 string,  
  value2 int,  
  value3 int)  
...  
WITH SERDEPROPERTIES ("hbase.columns.mapping"= ":key,a:b,a:c,d:e")  
TBLPROPERTIES("hbase.table.name" = "MY_TABLE");
```



Configuring Hive

- **Very little configuration is necessary to get started with Hive**
- **Minimum configuration identifies where to find the metastore**
 - If no configuration is provided a local Derby database is used
 - BigInsights installs pre-configured to use it's Derby server
 - Can be configured to use a wide variety of storage options (DB2, MySQL, Oracle, XML files, etc.)
- **Configuration file is located in:**
`$HIVE_HOME/conf/hive-site.xml`
- **Configuration Example :**

```
<property>  
  <name>hive.exec.scratchdir</name>  
  <value>/tmp/mydir</value>  
  <description>Scratch space for Hive jobs</description>  
</property>
```

Configuration Option Settings

Variable Name	Description
hive.ddl.output.format	The data format to use for DDL output (e.g. DESCRIBE table). One of "text" (for human readable text) or "json" (for a json object).
hive.exec.script.wrapper	Wrapper around any invocations to script operator e.g. if this is set to python, the script passed to the script operator will be invoked as python <script command>. If the value is null or not set, the script is invoked as <script command>.
hive.exec.scratchdir	This directory is used by hive to store the plans for different map/reduce stages for the query as well as to store the intermediate outputs of these stages.
hive.exec.submitviachild	Determines whether the map/reduce jobs should be submitted through a separate jvm in the non local mode.
hive.exec.script.maxerrs ize	Maximum number of serialization errors allowed in a user script invoked through TRANSFORM or MAP or REDUCE constructs.
hive.exec.compress.out put	Determines whether the output of the final map/reduce job in a query is compressed or not.
hive.default.fileformat	Default file format for CREATE TABLE statement. Options are TextFile, SequenceFile, RCFile, and Orc.

Conclusion

- **There are plenty of more interesting topics and features**

- Settings to control hive behavior
- Hints to control query execution strategy
- Built-in functions and user defined functions (and aggregates)
- Views
- Performance tuning
- Etc.

- **Limitations**

- Many limitations have been addressed in Hive 0.11 and 0.12
 - Added data type support, added window functions, better JDBC drivers, better Hive server, ...
- Still serious limitations
 - No subqueries
 - Limited HBase support
 - ...

Questions?

